Game Developers Conference 2011

Approximating Translucency for a Fast, Cheap and Convincing Subsurface Scattering Look



Colin Barré-Brisebois (speaker)
Marc Bouchard



Agenda

- Prelude Real-Time Demo
- 1. Translucency in Computer Graphics
- 2. Technique Details
- 3. Implementation Details
- Q & A



Fig. 1 – Real-Time Translucency in Frostbite 2



Real-Time Translucency Demo







Translucency in Computer Graphics

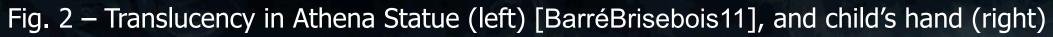


Translucency

The quality of allowing light to pass partially and diffusely inside media.









Translucency in Computer Graphics

- We rely heavily on BRDFs for describing local reflections
 - Simple and effective for opaque objects
- However, many objects in nature are (partly) translucent
 - Light transport also happens within the surface
 - BRDFs are not sufficient
- BSSRDFs allow for an even better simulation
 - But are usually more/too expensive
- In our case, we chose BSDFs (BRDF + BTDF),
 with some elements of BSSRDF

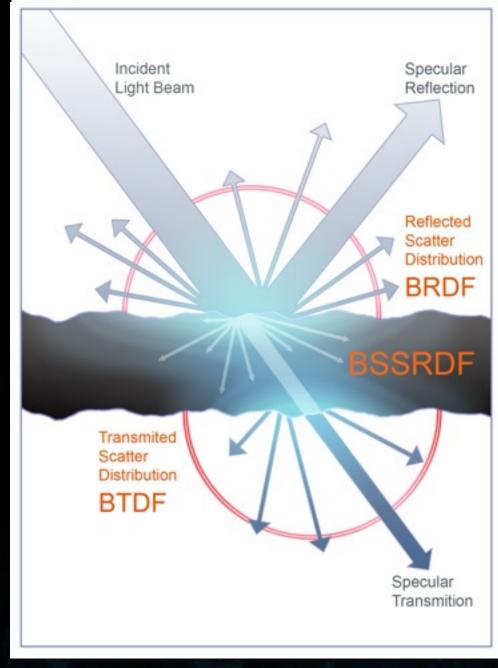


Fig. 3 – BRDF, BTDF and BSSRDF



The State of Translucency

Real-time translucency and derivatives come in different flavors:

- The more complex, but (relatively) expensive
 - [Chang08] Texture-space Importance Sampling
 - [Hable09] Texture-space Diffusion blurs, for skin / SSS
 - [Ki09] Shadowmap-based Translucency / SSS



Fig. 4 – Texture-space Importance Sampling [Chang08]



Fig. 5 – Texture-space Diffusion [Hable09]



The State of Translucency (cont.)

Real-time translucency and derivatives come in different

flavors:

- The simpler, but faster
 - [Sousa08] double-sided lighting & attenuation for foliage
- We want: fast like [Sousa08], and nice / complex results like the others ©

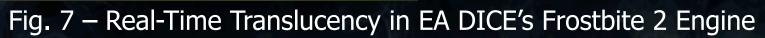


Fig. 6 – Foliage Translucency from Crytek's Crysis



And we got...







And we got... (cont.)





Technique Details



Overview

- We don't want to rely on additional depth maps and texture-space blurs
 - Requires more memory (i.e. for depth maps)
 - Requires significant computation (i.e. for texture blurs)
 - The previous are still feasible, but what if we could do without...
- In essence, for convincing translucency, the light traveling inside the shape:
 - Has to be influenced by the varying thickness of the object
 - Has to show some view & light-dependent diffusion/attenuation
- We only really need a simple representation of inner surface diffusion
 - Most users will be convinced even if not fully accurate!
 - Also, if the effect is cheap, we are free to use it everywhere!* ⊚



Overview (cont.)

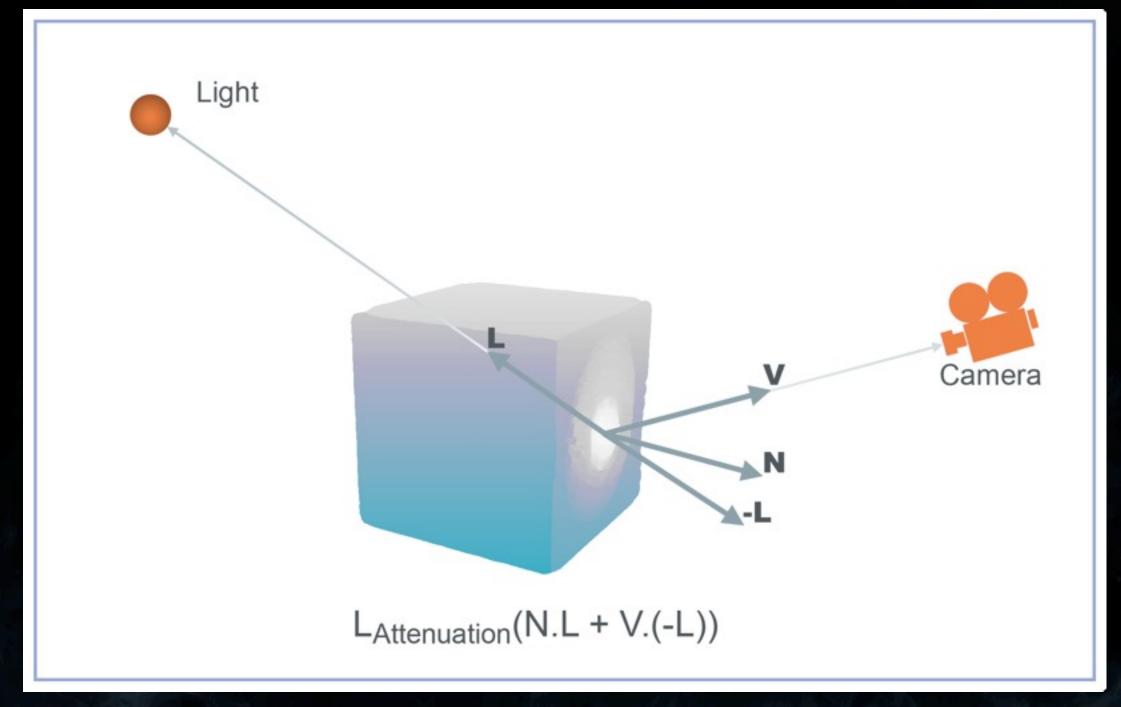
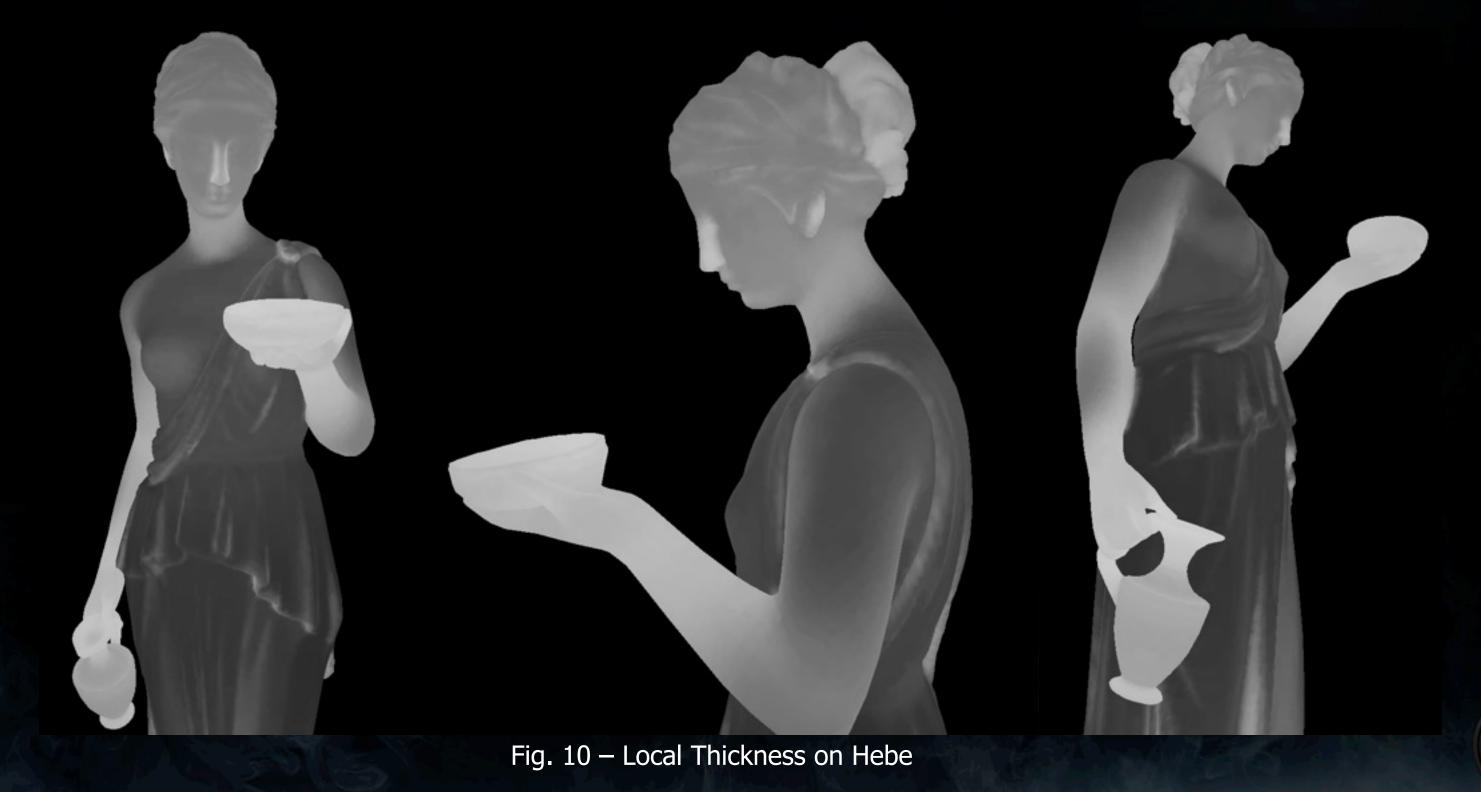


Fig. 9 – Direct and Translucency Lighting Vectors



Local Thickness





Computing Local Thickness

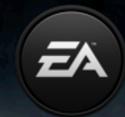
- We rely on ambient occlusion for computing this info:
 - 1. Invert the surface normals
 - 2. Render ambient occlusion
 - 3. Invert the colors and store in texture
- Can also be stored in vertex color, if tesselation allows
- Similar to [Sousa08], but streamlined for meshes of varying shapes and sizes.



Computing Local Thickness (cont.)



Fig. 11 – Local Thickness Texture for Hebe



What About Subsurface Scattering?

Even if not mathematically perfect, our technique gives an impression of SSS:

- Local thickness approximates light transport inside the shape
 - Different color for direct and indirect light gives convicing light color subsurface transfer
- View-oriented distortion and attenuation gives an organic result, breaking the uniformity



Fig. 13 – Our Technique, combined with Skin Shading



Implementation Details



Code

```
half3 vLTLight = vLight + vNormal * fLTDistortion;
half fLTDot = pow(saturate(dot(vEye, -vLTLight)), iLTPower) * fLTScale;
half3 fLT = fLightAttenuation * (fLTDot + fLTAmbient) * fLTThickness;
outColor.rgb += cDiffuseAlbedo * cLightDiffuse * fLT;
```

- Generates approx. 13 ALU instructions (based on platform)
 - More performance details in following slides
- Can precompute powers in order to get rid of the pow()

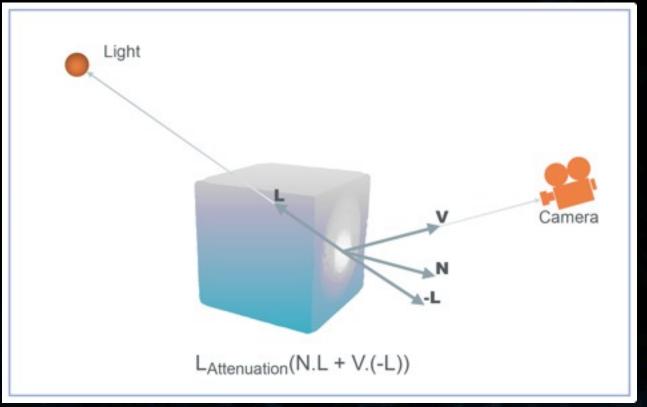


Fig. 9 – Direct and Translucency Lighting Vectors



Managing the Data at Runtime

At runtime, we have several parameters to manage, some perlight, some per-material:

- Per-light parameters are used during the deferred light pass
- Per-material parameters are stored in the g-buffer
 - Can also be stored in a separate buffer, if space-limited
 - Some parameters make more sense per-light, some per-material
 - This is very specific to your g-buffer setup
 - Might require some clever packing/unpacking
 - For packing more parameters on a per-material basis, instead of per-light
 - Using this technique with deferred shading definitely brings-the-thunder!



```
half3 vLTLight = vLight + vNormal * fLTDistortion;
half fLTDot = pow(saturate(dot(vEye, -vLTLight)), iLTPower) * fLTScale;
half3 fLT = fLightAttenuation * (fLTDot + fLTAmbient) * fLTThickness;
outColor.rgb += cDiffuseAlbedo * cLightDiffuse * fLT;
```

fLTAmbient

- Ambient value
- Visible from all angles
- Representing both front and back translucency that is always present
- Optimally, per-material

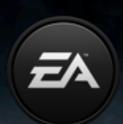


```
half3 vLTLight = vLight + vNormal * fLTDistortion;
half fLTDot = pow(saturate(dot(vEye, -vLTLight)), iLTPower) * fLTScale;
half3 fLT = fLightAttenuation * (fLTDot + fLTAmbient) * fLTThickness;
outColor.rgb += cDiffuseAlbedo * cLightDiffuse * fLT;
```

<u>iLTPower</u>

- Power value for direct translucency
- Breaks continuity, view-dependent
- Can be optimized with pre-computed powers
- Optimally, per-material

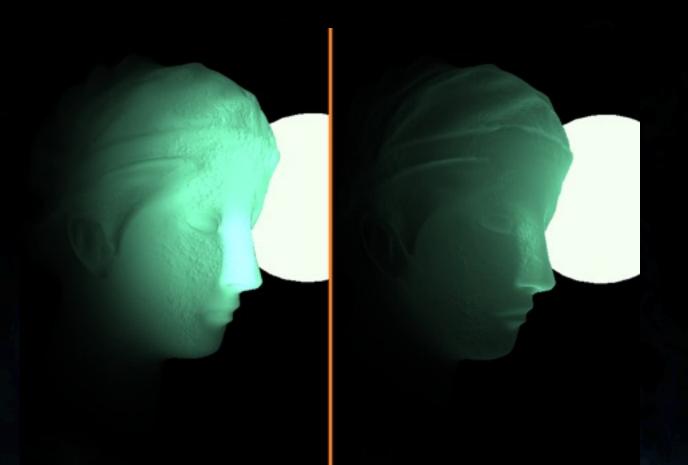




```
half3 vLTLight = vLight + vNormal * fLTDistortion;
half fLTDot = pow(saturate(dot(vEye, -vLTLight)), iLTPower) * fLTScale;
half3 fLT = fLightAttenuation * (fLTDot + fLTAmbient) * fLTThickness;
outColor.rgb += cDiffuseAlbedo * cLightDiffuse * fLT;
```

<u>fLTDistortion</u>

- Subsurface Distortion
- Shifts the surface normal
- Breaks continuity, view-dependent Allows for more organic, Fresnel-like
- Optimally, per-material





```
half3 vLTLight = vLight + vNormal * fLTDistortion;
half fLTDot = pow(saturate(dot(vEye, -vLTLight)), iLTPower) * fLTScale;
half3 fLT = fLightAttenuation * (fLTDot + fLTAmbient) * fLTThickness;
outColor.rgb += cDiffuseAlbedo * cLightDiffuse * fLT;
```

fLTThickness

- Pre-computed Local Thickness Map
- Used for both direct and indirect translucency
- Attenuates the computation where surface thickness varies
- Defined per-material

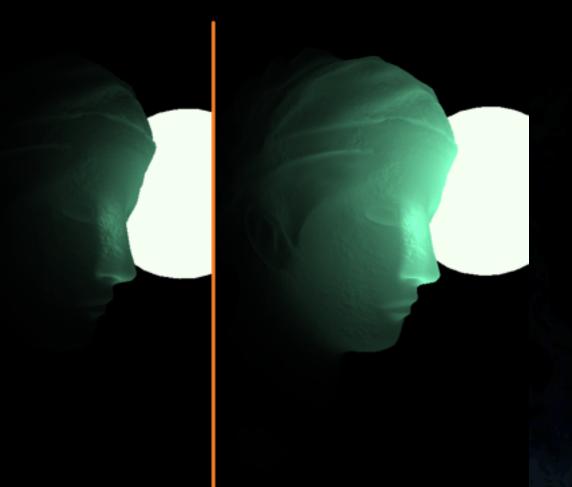




```
half3 vLTLight = vLight + vNormal * fLTDistortion;
half fLTDot = pow(saturate(dot(vEye, -vLTLight)), iLTPower) * fLTScale;
half3 fLT = fLightAttenuation * (fLTDot + fLTAmbient) * fLTThickness;
outColor.rgb += cDiffuseAlbedo * cLightDiffuse * fLT;
```

fLTScale

- Scale value
- Direct / Back translucency
- View-oriented
- Should be defined per-light. This makes it the central control point





All-Together





Deferred Shading G-Buffer Setup

Minimally, translucency can be stored in the g-buffer as a single greyscale value:

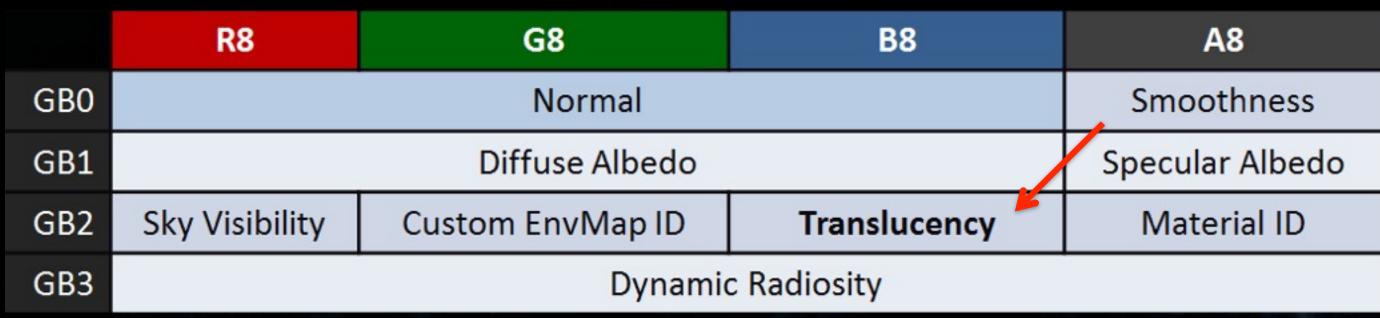


Fig. 20 – Our G-Buffer, with greyscale Translucency

Based on your game, this can be enough. The color will then only originate from the light sources (and also diffuse albedo).

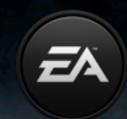


Deferred Shading G-Buffer Setup (cont.)

All objects here are relying on a greyscale value for translucency \rightarrow



Fig. 1 – Real-Time Translucency in Frostbite 2



Deferred Shading G-Buffer Setup (cont.)

Better results will be achieved if translucency is a color (here, with some packing):

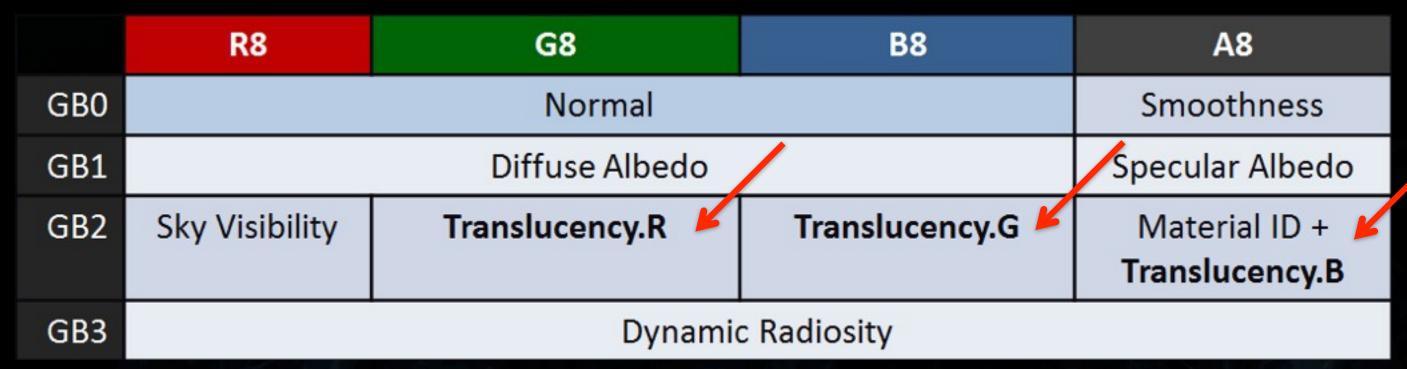


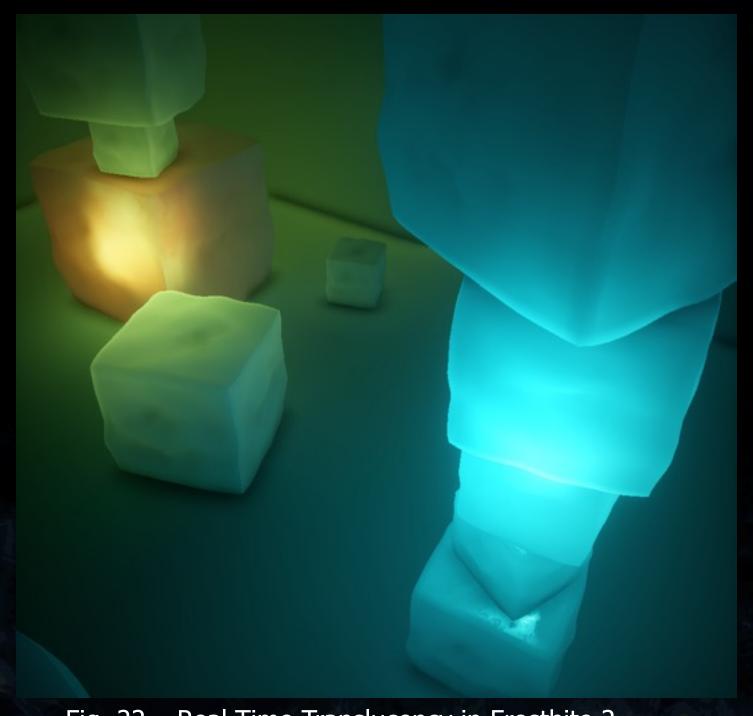
Fig. 21 – Our G-Buffer, with coloured Translucency (packed/offset)

This translucency color, representing our inner surface color diffusion, will be combined to the light color and the material's diffuse albedo.



Deferred Shading G-Buffer Setup (cont.)

Green Light
White Albedo
Red Translucency ->



<- Blue Light
White Albedo &
Translucency



Performance

	XBOX 360	PLAYSTATION 3	PC (DX11)
Full-Screen Coverage	0.6 ms	1.0 ms	0.03 ms
Instructions	13	17	12

- PS3: Takes advantage of our light-tile rendering on SPUs
 - -See Christina Coffin's "SPU Deferred Shading for BF3 on PS3"
- DX11: Supported in our Compute Shader solution
 - —See Johan Andersson's "DX11 rendering in Battlefield 3" talk for more DX11-related details
- PC: AMD Radeon 6970



Caveats

- Doesn't take all concavities into account
- Technique is optimal for convex hulls
- Doesn't work with morphing/animated objects
 - Alternative: Though camera dependent, several cases could work with a real-time thickness approximat [Oat08]
 - Alternative: Could also use an hybrid dynamic AO computation, with inverted normals

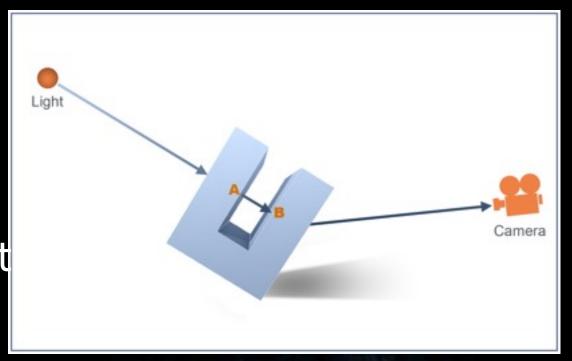
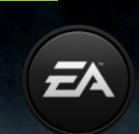


Fig. 23 – Concave Hull



Summary and Q&A

- We presented an artist-friendly, fast and scalable real-time approximation of light transport in translucent homogenous media:
 - -Improves games visuals by adding a new dimension, with light traveling inside shapes
 - -Has a scalable and reasonable impact on runtime
 - Provides convincing results even if not mathematically perfect
- This technique is also published in GPU Pro 2, released this week at GDC. Check out the book!



GPU Pro²

Questions?



Email: cbbrisebois@ea.com

mbouchard@ea.com

Blog: http://colinbarrebrisebois.com

Twitter: @ZigguratVertigo

Battlefield 3 & Frostbite 2 talks at GDC'11:

Mon 1:45	DX11 Rendering in Battlefield 3	Johan Andersson
Wed 10:30	SPU-based Deferred Shading in Battlefield 3 for PlayStation 3	Christina Coffin
Wed 3:00	Culling the Battlefield: Data Oriented Design in Practice	Daniel Collin
Thu 1:30	Lighting You Up in Battlefield 3	Kenny Magnusson
Fri 4:05	Approximating Translucency for a Fast, Cheap & Convincing Subsurface Scattering Look	Colin Barré-Brisebois



For more DICE talks: http://publications.dice.se

Special Thanks

- Sergei Savchenko
- Johan Andersson (@repi)
- Christina Coffin (@christinacoffin)
- Halldor Fannar
- Joakim Svärling
- Stephen Hill (@self_shadow)
- Frederic O'Reilly
- John White
- Wessam Bahnassi
- Carsten Dachsbacher
- Daniel Collin (@daniel_collin)
- Torbjörn Malmer
- Kenny Magnusson
- Dominik Bauset
- Sandra Jensen



- Mohannad Al-Khatib (@psychodesigns)
- Colin Boswell (@bozz)



References

[BarréBrisebois11] Barré-Brisebois, Colin and Bouchard, Marc."Real-Time Approximation of Light Transport

in Translucent Homogenous Media", GPU Pro 2, Wolfgang Engel, Ed. Charles River

Media, 2011.

[Chang08] Chang, Chih-Wen, Lin, Wen-Chieh, Ho, Tan-Chi, Huang, Tsung-Shian and Chuang,

Jung-Hong. "Real-Time Translucent Rendering Using GPU-based Texture Space

Importance Sampling," Computer Graphics Forum (Eurographics 2008), Vol. 27,

No. 2, 2008, pp 517-526.

[Hable09] Hable, John, Borshukov, George and Hejl, Jim. "Fast Skin Shading," *ShaderX7: Advanced*

Rendering Techniques, Wolfgang Engel, Ed., Charles River Media, 2009: pp. 161-173.

[Ki09] Ki, Hyunwoo. "Real-time Subsurface Scattering Using Shadow Maps," ShaderX7:

Advanced Rendering Techniques, Wolfgang Engel, Ed., Charles River Media, 2009:

pp. 467-478.

[Oat08] Oat, Christopher and Scheuermann, Thorsten. "Computing Per-Pixel Object Thickness in

a Single Render Pass," ShaderX6: Advanced Rendering Techniques, Wolfgang Engel,

Ed., Charles River Media, 2008: pp. 57-62.

[Sousa08] Sousa, Tiago. "Vegetation Procedural Animation and Shading in Crysis," GPU Gems 3,

Hubert Nguyen, Ed., Addison-Wesley, 2008: pp. 373-385.



